

EDA

Student: Femke Bakker (13264443)

- Goal of the research: classify documents using LLMs
- The data consists of documents and their labels (classes). Besides the text from the pdfs there are no other features that will be used for the classification
- The EDA is performed on 20% of the dataset, because the dataset is very large (33000 docs).
- First the data is loaded in and some cleaning of the tokens is done.
- Missing values, documents where no text was extracted from the PDFs, are already removed.

Link to github: <https://github.com/Amsterdam-Internships/document-classification-using-large-language-models/tree/main>

```
import sys
sys.path.append('../scripts/')

import pandas as pd
from load_data import load_txt_files
import nltk
from nltk.tokenize import word_tokenize
import matplotlib.pyplot as plt

txtfile_paths = pd.read_csv("../data/txtfile_paths.csv")
all_txt_df = load_txt_files(txtfile_paths, ['complete'])
df = all_txt_df.copy()

df['tokens'] = df['text'].apply(word_tokenize)

def calculate_length(token_list):
    return len(token_list)

df['token_length'] = df['tokens'].apply(calculate_length)
SAVED_DF = df.copy()

import numpy as np

missing_text = df[df['text'].str.len() < 5]
print(f"There are {len(missing_text)} documents that have less than 5 characters")

print(f"Removing the documents that could not be extracted or have less than 5 characters, leaves us with {len(df)-len(missing_text)} (out of {len(df)})")
df = df[df['text'].str.len() > 5]
```

There are 11 documents that have less than 5 characters
Removing the documents that could not be extracted or have less than 5 characters, leaves us with 33117 (out of 33128)

```
from nltk.corpus import stopwords
import string
import nltk
nltk.download('stopwords')

# remove stopwords
top_df = df.copy()

def remove_stopwords(tokens):
    stop_words = set(stopwords.words('dutch'))
    tokens_without_stopwords = [word for word in tokens if
word.lower() not in stop_words]
    tokens_without_punctuation = [word for word in
tokens_without_stopwords if word not in string.punctuation and
len(word)>1]
    return tokens_without_punctuation

df['cleaned_tokens'] = df['tokens'].apply(remove_stopwords)
df['clean_text'] = df['cleaned_tokens'].apply(lambda x: ' '.join(x))

[nltk_data] Downloading package stopwords to
[nltk_data] /home/azureuser/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

Token & class distribution

Below are the tables shown for the complete dataset and the classes.

We can see that the token counts within the classes and between the classes are varying. This could create difficulties when converting the docs into suitable input for the LLMs. The LLMs have a max token limit, meaning that whole documents cannot be given as input and thus the docs need to be represented in another way. Furthermore, we can see the distribution of the amount of docs per class is some what even, with begroting being an outlier, with a lot of docs. Actualiteit and Factsheets are a little underrepresented compared to the other classes.

Note: the minimum values are very small. This is likely a result of mistakes made during converting PDF to OCR. Further expection is needed to remove bad files.

```
display(df['token_length'].describe())
class_describe = df.groupby('label')['token_length'].describe()
display(class_describe)

classes = list(set(df['label']))
counts = list(class_describe['count'])
```

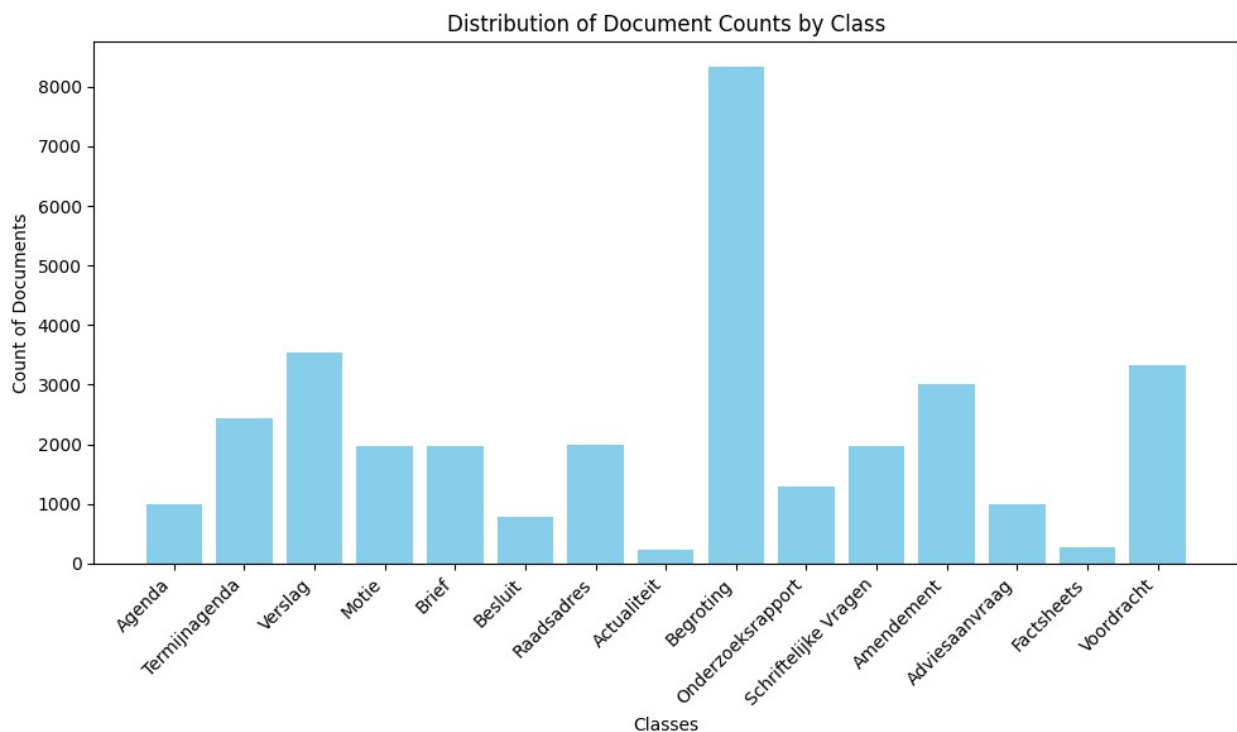
```
# Plotting the distribution
plt.figure(figsize=(10, 6))
plt.bar(classes, counts, color='skyblue')
plt.xlabel('Classes')
plt.ylabel('Count of Documents')
plt.title('Distribution of Document Counts by Class')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

```
count      33117.000000
mean        2691.847933
std         13409.818780
min           3.000000
25%         291.000000
50%         533.000000
75%        1225.000000
max        275597.000000
Name: token_length, dtype: float64
```

	count	mean	std	min
25% \ label				
Actualiteit 236.00	996.0	696.773092	3462.293848	72.0
Adviesaanvraag 557.00	2442.0	1644.871417	2993.302964	57.0
Agenda 325.00	3537.0	1048.904156	4244.783305	20.0
Amendement 236.00	1969.0	2658.493652	7044.470410	30.0
Begroting 248.00	1967.0	13160.190646	46780.793626	40.0
Besluit 168.50	775.0	986.649032	1749.006500	70.0
Brief 734.00	1995.0	1764.259649	1728.612069	3.0
Factsheets 1154.00	234.0	6008.987179	14742.837685	112.0
Motie 234.00	8336.0	521.707893	1129.571368	105.0
Onderzoeksrapport 5896.25	1286.0	15329.244168	18086.586572	233.0
Raadsadres 358.50	1975.0	1072.450633	1666.327342	34.0
Schriftelijke Vragen 871.75	3004.0	1742.592543	5832.719196	26.0
Termijnagenda	996.0	545.188755	346.153706	48.0

257.00				
Verslag	273.0	36233.520147	12422.807218	2398.0
29160.00				
Voordracht	3332.0	619.479892	497.505838	121.0
368.00				

	50%	75%	max
label			
Actualiteit	344.5	592.0	103700.0
Adviesaanvraag	955.0	1672.0	77769.0
Agenda	616.0	939.0	139648.0
Amendement	359.0	1184.0	62304.0
Begroting	385.0	2324.0	247184.0
Besluit	439.0	1216.0	26523.0
Brief	1269.0	2291.5	32957.0
Factsheets	2753.5	5462.0	171297.0
Motie	292.5	402.0	36091.0
Onderzoeksrapport	10999.0	19127.5	275597.0
Raadsadres	627.0	1104.0	24021.0
Schriftelijke Vragen	1209.5	1675.0	181725.0
Termijnagenda	486.0	761.5	3338.0
Verslag	38222.0	45151.0	59938.0
Voordracht	468.0	673.5	5390.0



```
import matplotlib.pyplot as plt
```

```

plt.figure(figsize=(4, 3))
plt.boxplot(df['token_length'], vert=False)
plt.title('Boxplot of Token Length for the whole dataset')
plt.xlabel('Token Count')
plt.show()

grouped = df.groupby('label')

# Create subplots
fig, axs = plt.subplots(4, 4, figsize=(10, 7.5), sharey=True)

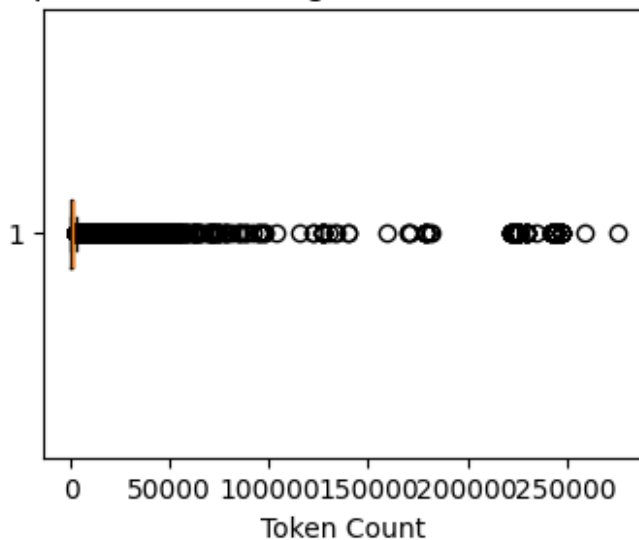
# Iterate over groups and plot boxplots
for i, (label, group) in enumerate(grouped):
    row = i // 4
    col = i % 4
    axs[row, col].boxplot(group['token_length'], vert=False)
    axs[row, col].set_title(f'Class {label}')
    axs[row, col].set_xlabel('Token Count')
    axs[row, col].set_ylabel('Class')

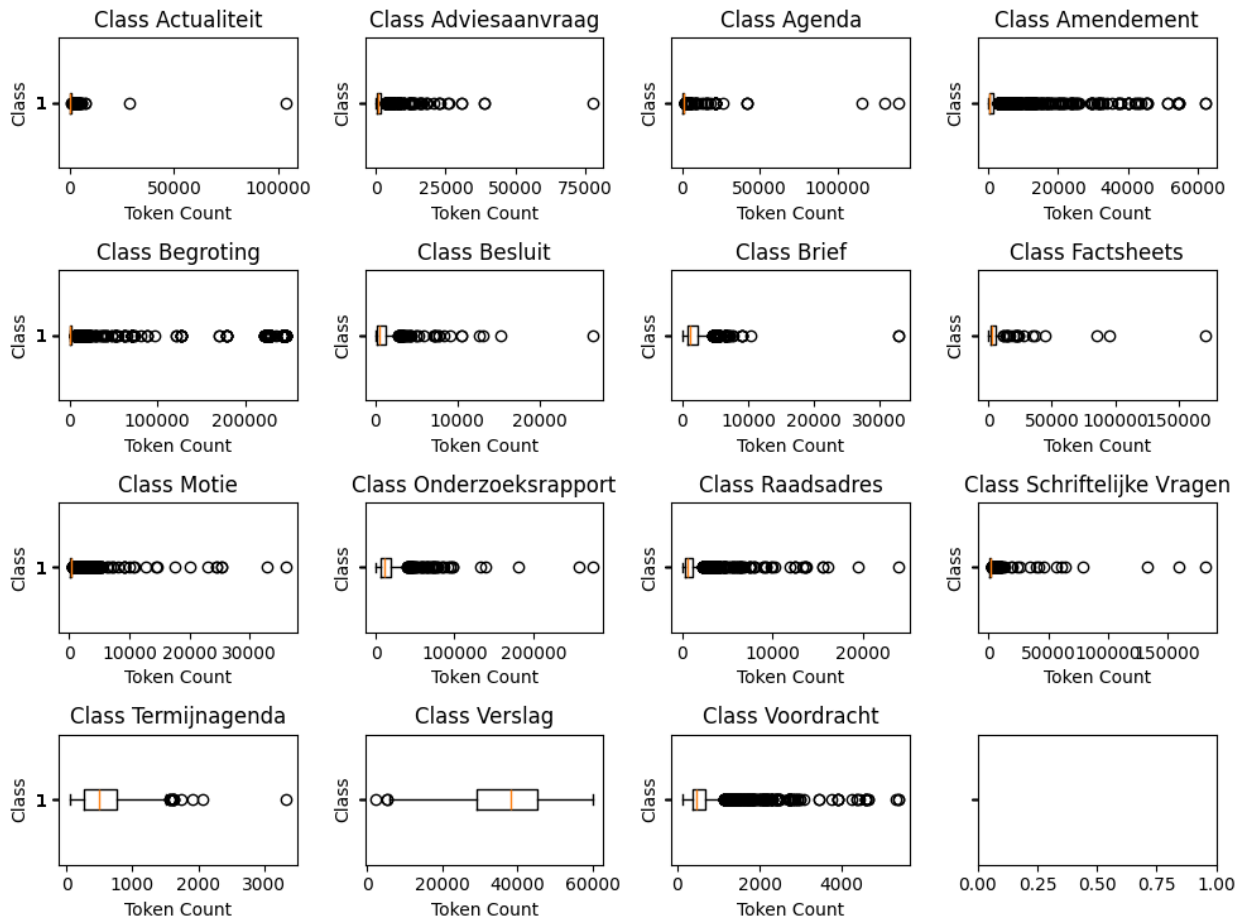
# Hide empty subplots if any
for i in range(len(grouped), 3*5):
    row = i // 5
    col = i % 5
    axs[row, col].axis('off')

plt.tight_layout()
plt.show()

```

Boxplot of Token Length for the whole dataset





Based on the plots above we can see that there are classes that don't have (many) outliers in token counts, such as termijnagenda, verslag and brief. On the other hand there are also classes that have many outliers in token count, such as voordracht, begroting and moties. It's noticeable that the outliers are on the right side, meaning that the outliers have a high token count.

Top words for each class

The goal is to find words that are identifiers for the classes. We suspect that the name of the class, for example 'agenda', 'motie', 'begroting', will be strong indicators for a class. However, we also know that these words are used in documents that don't belong to that class. The goal of this part of the analysis is to find how severe this is.

1. select top 50 words with highest TF-IDF mean for each class.
2. Then select the unique words for each class. Words that do not occur in the top-50 of other classes.
3. Check if other class names are named in docs of a class.

```
import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
```

```

tfidf_vectorizer = TfidfVectorizer()
tfidf_matrix = tfidf_vectorizer.fit_transform(df['clean_text'])
tfidf_df = pd.DataFrame(tfidf_matrix.toarray(),
columns=tfidf_vectorizer.get_feature_names_out())
tfidf_df['label'] = df['label']

# Calculate mean TF-IDF score for each word across all documents in
each class
mean_tfidf_by_class = tfidf_df.groupby('label').mean()

# Get top 50 words with highest mean TF-IDF score for each class
top_50_words = {}
for class_name, class_tfidf in mean_tfidf_by_class.iterrows():
    top_50_words[class_name] = class_tfidf.nlargest(50).index.tolist()

unique_words_dict = {}
for class_name in top_50_words.keys():
    top_words = top_50_words[class_name]

    all_other_words = []
    for other_name in top_50_words.keys():
        if class_name != other_name:
            all_other_words.extend(top_50_words[other_name])
    all_other_words = set(all_other_words)

    unique_words = [word for word in top_words if word not in
all_other_words]
    print(f'{class_name}: ---- {unique_words}')

```

```

Actualiteit: ---- ['actualiteit', 'spoedeisendheid', 'reden',
'raadsactualiteit', 'behandeling', 'maart', 'poot', 'supplement',
'wassink', 'februari', 'indieners']
Adviesaanvraag: ---- ['stadsdeelcommissie', 'adviesaanvraag', 'art',
'concept', 'stadsdelen', 'staf', 'invullen', 'adviezen', '2605',
'81484', 'z17', 'sdc', 'verordening', 'directie', 'kaders', 'kader',
'meegestuurd', 'weekstart', 'conceptadvies', 'procesbegeleider',
'formuleer', 'vraagt', 'optioneel']
Agenda: ---- ['dient', 'aanvang', 'bd2013']
Amendement: ---- ['pvda', 'stadsdeelraad', 'sp', 'dhr', 'groenlinks',
'deelraad', 'bijeen', 'bnw81', 'cda', 'dag']
Begroting: ---- ['2010', '2009', '2008', 'miljoen', '2007',
'december', 'lasten', 'baten', 'reserve', 'euro', 'reserves',
'kosten', 'middelen', 'saldo', 'behandelen', 'bedrag']
Besluit: ---- ['algemeen', 'ab', 'int', 'bestemmingsplan', 'decos',
'nieuwwest', 'ivar', 'baâdoud', '1064', '14020', 'wink', 'sw', '2003',
'plein', 'stadsdeelsecretaris', 'bestuurscommissies', 'ca',
'bezoekadres', 'vast', 'ontwerpbestemmingsplan', 'afschrift', 'gelet']
Brief: ---- ['routebeschrijving', 'vindt', 'portefeuillehouder',

```

'2024', 'geïnformeerd', 'maatregelen', 'ontwikkeling', 'portefeuille', 'verschillende', 'vit', 'daarnaast']

Factsheets: ---- ['aandeel', 'bron', '65', 'minima', 'huishoudens', 'minimahuishoudens', 'gemiddelde', 'the', 'zuidoost', 'buurten', 'gemiddeld', 'and', 'inkomen', 'factsheet', 'cbs', '18', 'ouderen', 'cijfers', 'ois', 'leerlingen']

Motie: ---- ['status', 'verzoekt', 'ondergetekenden', 'verworpen', 'amsterdamse']

Onderzoeksrapport: ---- ['es', 'ie', 're', 'ed', 'el', 'ke', 'pe', 'et', 'tussen', 'figuur', 'mn', 'he', 'le', 'scholen', 'be']

Raadsadres: ---- ['kenmerk', '202', '020', 'mail', 'www', 'graag', 'bijlage', '552', 'komen', 'alleen', 'verzonden', 'binnenstad', 'retouradres']

Schriftelijke Vragen: ---- ['antwoord', 'schriftelijke', 'politie', 'beantwoording', 'vorenstaande', 'welke', 'grond', 'nee', 'toelichting', 'vragensteller', 'waarom', 'reglement', 'bekend', 'akkoord', '45', 'neng', 'augustus']

Termijnagenda: ---- ['termijnagenda', 'bd2015', 'bd2012', 'actualiteiten', '15', 'cie']

Verslag: ---- ['heel', 'raadsnotulen', 'woord', 'even', 'denk', 'natuurlijk', 'goed', 'echt', 'vind', 'zeggen', 'eigenlijk', 'gewoon', 'geeft', 'willen', 'misschien', 'dank', 'stemming', 'gezegd', 'zegt', 'zitten', 'weet']

Voordracht: ---- ['pdf', 'nvt', 'gegenereerd', 'vl', 'kennisneming', 'ad2023', 'naam', 'bespreking', 'ad2022', 'ad2021', 'gemeentewet', 'geheimhouding', 'raadscommissies', 'toezegging', 'tijdelijke', 'indienend', 'uitgenodigde', 'afgedaan', 'mailadres', 'behandelend', 'treft', 'telefoonnummer', 'extern', 'grondslag', 'ambtenaar', 'algemene', 'inzage', 'wettelijke', 'achtergrond', 'bestuurlijke', 'gevraagd', 'uitkomsten']

get dict with all tokens (including duplicates) for each class

```
tokens_class = dict()
for class_name in set(df['label']):
    all_tokens = list(df.loc[df['label']==class_name]
['cleaned_tokens'].values)
    all_tokens = [token for sublist in all_tokens for token in
sublist]
    tokens_class[class_name] = all_tokens
```

For each class count how many times the class names are in the tokens

```
df_col = list(tokens_class.keys())
df_col.append("Total Doc in Class")
class_count_df = pd.DataFrame(columns=df_col)
```

```
for class_name in tokens_class.keys():
    tokens = tokens_class[class_name]
    counts = {label: tokens.count(label.lower()) for label in
```



```

tokens_class.keys()}
    counts['Total Doc in Class'] =
len(df.loc[df['label']==class_name])
    class_count_df.loc[len(class_count_df)] = counts

class_count_df.set_index(pd.Index(list(tokens_class.keys()))),
inplace=True)

# high the df to make result more interpretable
def highlight_max_and_second_highest_except_total(s):
    max_val = s[:-1].max()
    second_max_val = s[:-1].nlargest(2).iloc[-1] # Get the second
largest value
    is_max = s[:-1] == max_val
    is_max['Total Doc in Class'] = False
    is_second_max = s[:-1] == second_max_val
    is_second_max['Total Doc in Class'] = False

    styles = ['background-color: green' if v else '' for v in is_max]
    styles_second = ['background-color: orange' if v else '' for v in
is_second_max]
    combined_styles = [f'{styles[i]}; {styles_second[i]}' for i in
range(len(styles))]
    return combined_styles

styled_df =
class_count_df.style.apply(highlight_max_and_second_highest_except_tot
al, axis=1)

display(styled_df)

<pandas.io.formats.style.Styler at 0x7f46c7bd1070>

```

The green marked values are the highest count in that row and the orange ones are the second highest. We want the class_name to be named a lot in the class and very few in the other classes, that means that the class name is an identifier for that class.

Notes:

- schriftelijke vragen is never named in the docs. Which is understandable, since schriftelijke vragen are questions about very different subjects
- the distinction between agenda and termijnagenda might be hard, based on the count of agenda and termijnagenda in them. They both name agenda a lot.
- Verslag might be confused with Motie, since motie is named very often.

Overall there is quite some overlap between the classes, in regards of naming the class names.

Similarity within class

If the docs within a class are similar, it will be easier to classify them, because then patterns are easier to identify by the model. The average cosine similarity between docs of a class is calculated. Since the texts need to be shortened to give as input, the similarity of the first 1000 tokens of each doc is also compared.

We can see in the tables below the similarity of the documents within the classes. It shows that the documents are not similar to each other, since the scores are low. Verslag is the class with the highest similarity between the documents. Additionally, the documents are more similar if the whole document is compared, than when the first 1000 tokens of the docs are compared. This could be a problem, or at least worse the performance if the docs are represented by the first N tokens. This could mean that other representation methods such as summarizing might be a better fit to represent the docs.

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
import math

def similarity_within_class(input_df, text_len):
    work_df = input_df.copy()
    within_sim = pd.DataFrame(columns=['label', 'average_sim',
    'num_pairs_exceeding_threshold'])

    # if not complete
    if text_len == 'complete':
        text_column = 'text'

    else:
        work_df['short_text'] = work_df['tokens'].apply(lambda tokens:
        ' '.join(tokens[:text_len]))
        text_column = 'short_text'

    for category in set(work_df['label']):
        subdf = work_df.loc[work_df['label']==category]

        # get tf-idf score
        tfidf_vectorizer = TfidfVectorizer()
        tfidf_matrix =
        tfidf_vectorizer.fit_transform(subdf[text_column])

        # calculate cosine similarity
        cosine_similarities = cosine_similarity(tfidf_matrix,
        tfidf_matrix)

        # take average sim
        average_similarity = round(np.mean(cosine_similarities), 2)
```

```

        # calculate how many pairs succeed threshold
        total_possible_pairs = len(cosine_similarities) *
(len(cosine_similarities))
        num_pairs_exceeding_threshold =
round(np.sum(cosine_similarities > 0.5) / total_possible_pairs, 2)

        # add to df
        within_sim.loc[len(within_sim)] = {'label':category,
'average_sim': average_similarity,
'num_pairs_exceeding_threshold':num_pairs_exceeding_threshold}

        within_sim = within_sim.sort_values(by='average_sim',
ascending=False)
        return within_sim

# sim_within_df_500tokens = similarity_within_class(df, 500)
# display(sim_within_df_500tokens)

sim_within_df_1000tokens = similarity_within_class(df, 1000)
display(sim_within_df_1000tokens)

sim_within_df_cleantext = similarity_within_class(df, 'complete')
display(sim_within_df_cleantext)

```

	label	average_sim	num_pairs_exceeding_threshold
2	Verslag	0.47	0.32
12	Adviesaanvraag	0.28	0.01
4	Brief	0.26	0.00
10	Schriftelijke Vragen	0.25	0.00
0	Agenda	0.24	0.04
5	Besluit	0.20	0.01
1	Termijnagenda	0.19	0.05
7	Actualiteit	0.18	0.01
8	Begroting	0.18	0.00
9	Onderzoeksrapport	0.18	0.01
14	Voordracht	0.18	0.00
6	Raadsadres	0.17	0.00
13	Factsheets	0.17	0.02
11	Amendement	0.16	0.00
3	Motie	0.14	0.00

	label	average_sim	num_pairs_exceeding_threshold
2	Verslag	0.91	1.00
9	Onderzoeksrapport	0.48	0.47
12	Adviesaanvraag	0.34	0.05
4	Brief	0.33	0.05
8	Begroting	0.30	0.07
10	Schriftelijke Vragen	0.29	0.00
0	Agenda	0.26	0.05
13	Factsheets	0.25	0.05

5	Besluit	0.22	0.01
11	Amendement	0.22	0.03
1	Termijnagenda	0.19	0.05
6	Raadsadres	0.19	0.00
7	Actualiteit	0.19	0.01
14	Voordracht	0.18	0.00
3	Motie	0.15	0.00